

Wine Quality Prediction



Bharat Machine Learning Internship

Objective : To Predict Quality of Wine Using Linear Regression

Author : Sonar Daxita Ramakantbhai

Method Use : Linear Regression

Dataset : <https://www.kaggle.com/datasets/yasserh/wine-quality-dataset?resource=download>

Steps We are Followed

- 1) Reading the dataset
- 2) Checking and cleaning the data
- 3) Visualize the data
- 4) predict the wine Quality
- 5) Draw the conclusions

import the required library set

```
import pandas as pd
import numpy as np
import matplotlib as plt
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
import sklearn
from sklearn.model_selection import train_test_split
#supress warning
import warnings
warnings.filterwarnings("ignore")
```

Read the data

```
data_wine=pd.read_csv("wine.csv")
print("data has been successfully import..")
```

data has been successfully import..

Set directory

```
import os
os.getcwd()

'C:\\Users\\Admin\\Bharat_intern'

os.chdir("E:\\Bharat_intern\\Task_2")
os.getcwd()

'E:\\Bharat_intern\\Task_2'
```

#check first five observations

```
data_wine.head()
```

	fixed acidity chlorides \	volatile acidity	citric acid	residual sugar
0	7.4	0.70	0.00	1.9
0.076				
1	7.8	0.88	0.00	2.6
0.098				
2	7.8	0.76	0.04	2.3
0.092				
3	11.2	0.28	0.56	1.9
0.075				
4	7.4	0.70	0.00	1.9
0.076				

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	11.0	34.0	0.9978	3.51	0.56
1	25.0	67.0	0.9968	3.20	0.68
2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56

	alcohol	quality	Id
0	9.4	5	0
1	9.8	5	1
2	9.8	5	2
3	9.8	6	3
4	9.4	5	4

#check last five observations

data_wine.tail()

	fixed acidity	volatile acidity	citric acid	residual sugar
1138	6.3	0.510	0.13	2.3
1139	6.8	0.620	0.08	1.9
1140	6.2	0.600	0.08	2.0
1141	5.9	0.550	0.10	2.2
1142	5.9	0.645	0.12	2.0

	free sulfur dioxide	total sulfur dioxide	density	pH
1138	29.0	40.0	0.99574	3.42
1139	28.0	38.0	0.99651	3.42
1140	32.0	44.0	0.99490	3.45
1141	39.0	51.0	0.99512	3.52
1142	32.0	44.0	0.99547	3.57

	alcohol	quality	Id
1138	11.0	6	1592
1139	9.5	6	1593
1140	10.5	5	1594
1141	11.2	6	1595
1142	10.2	5	1597

#Checking the shape of data

data_wine.shape

(1143, 13)

#checking info of data

data_wine.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1143 entries, 0 to 1142
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1143 non-null   float64
1   volatile acidity       1143 non-null   float64
2   citric acid            1143 non-null   float64
3   residual sugar         1143 non-null   float64
4   chlorides              1143 non-null   float64
5   free sulfur dioxide    1143 non-null   float64
6   total sulfur dioxide   1143 non-null   float64
7   density                1143 non-null   float64
8   pH                    1143 non-null   float64
9   sulphates              1143 non-null   float64
10  alcohol                1143 non-null   float64
11  quality                1143 non-null   int64
12  Id                     1143 non-null   int64
dtypes: float64(11), int64(2)
memory usage: 116.2 KB

```

#statistical summary of the data
data_wine.describe()

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	1143.000000	1143.000000	1143.000000	1143.000000
mean	8.311111	0.531339	0.268364	2.532152
std	1.747595	0.179633	0.196686	1.355917
min	4.600000	0.120000	0.000000	0.900000
25%	7.100000	0.392500	0.090000	1.900000
50%	7.900000	0.520000	0.250000	2.200000
75%	9.100000	0.640000	0.420000	2.600000
max	15.900000	1.580000	1.000000	15.500000

	chlorides	free sulfur dioxide	total sulfur dioxide
density \			
count	1143.000000	1143.000000	1143.000000
1143.000000			
mean	0.086933	15.615486	45.914698
0.996730			
std	0.047267	10.250486	32.782130
0.001925			
min	0.012000	1.000000	6.000000
0.990070			
25%	0.070000	7.000000	21.000000
0.995570			
50%	0.079000	13.000000	37.000000
0.996680			
75%	0.090000	21.000000	61.000000
0.997845			

max	0.611000	68.000000	289.000000		
1.003690					
	pH	sulphates	alcohol	quality	Id
count	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000
mean	3.311015	0.657708	10.442111	5.657043	804.969379
std	0.156664	0.170399	1.082196	0.805824	463.997116
min	2.740000	0.330000	8.400000	3.000000	0.000000
25%	3.205000	0.550000	9.500000	5.000000	411.000000
50%	3.310000	0.620000	10.200000	6.000000	794.000000
75%	3.400000	0.730000	11.100000	6.000000	1209.500000
max	4.010000	2.000000	14.900000	8.000000	1597.000000

#checking the null value of our data

`data_wine.isnull().sum()`

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                0
sulphates          0
alcohol            0
quality            0
Id                0
dtype: int64
```

Exploratory Data Analysis

Data Visualization

```
for col in data_wine:
    print(data_wine[col].value_counts(ascending=False), '\n\n\n')
```

```
7.2    43
7.1    41
7.0    40
7.8    40
```

7.5	37
	..
4.6	1
13.7	1
13.4	1
13.5	1
12.2	1

Name: fixed acidity, Length: 91, dtype: int64

0.600	32
0.500	32
0.430	31
0.390	29
0.580	28
	..
1.035	1
0.565	1
0.865	1
0.965	1
0.160	1

Name: volatile acidity, Length: 135, dtype: int64

0.00	99
0.49	47
0.24	42
0.02	35
0.01	26
	..
0.61	1
0.72	1
1.00	1
0.75	1
0.62	1

Name: citric acid, Length: 77, dtype: int64

2.00	107
2.10	103
1.80	92
2.20	88
1.90	80
	...
7.30	1
7.20	1

2.95	1
3.65	1
4.40	1

Name: residual sugar, Length: 80, dtype: int64

0.080	48
0.077	41
0.074	38
0.084	38
0.078	36

	..
0.222	1
0.422	1
0.034	1
0.387	1
0.230	1

Name: chlorides, Length: 131, dtype: int64

6.0	99
5.0	80
12.0	58
10.0	52
15.0	51
7.0	51
9.0	48
16.0	47
8.0	45
17.0	40
11.0	39
13.0	39
14.0	38
18.0	37
3.0	33
19.0	32
4.0	31
21.0	30
23.0	23
26.0	21
24.0	21
27.0	21
25.0	20
20.0	18
32.0	18
31.0	16
28.0	15
29.0	14

22.0	12
34.0	12
30.0	10
36.0	9
35.0	9
33.0	8
38.0	8
41.0	5
48.0	4
1.0	3
40.0	3
42.0	3
43.0	2
45.0	2
52.0	2
51.0	2
37.5	2
68.0	2
39.0	2
46.0	1
53.0	1
40.5	1
55.0	1
37.0	1
66.0	1

Name: free sulfur dioxide, dtype: int64

28.0	36
15.0	28
14.0	27
20.0	27
18.0	26
	..
114.0	1
135.0	1
129.0	1
165.0	1
151.0	1

Name: total sulfur dioxide, Length: 138, dtype: int64

0.99760	27
0.99720	25
0.99680	22
0.99940	22
0.99640	21
	..

0.99438	1
0.99634	1
0.99426	1
0.99747	1
0.99651	1

Name: density, Length: 388, dtype: int64

3.30	41
3.36	40
3.38	38
3.39	37
3.26	33
..	
2.86	1
2.95	1
2.74	1
3.75	1
2.90	1

Name: pH, Length: 87, dtype: int64

0.60	53
0.62	50
0.56	47
0.54	46
0.57	42
..	
1.61	1
1.31	1
0.33	1
1.56	1
1.01	1

Name: sulphates, Length: 89, dtype: int64

9.500000	92
9.400000	72
9.800000	57
9.200000	50
10.000000	49
..	
11.950000	1
9.950000	1
9.233333	1
9.250000	1

```
10.550000      1
Name: alcohol, Length: 61, dtype: int64
```

```
5      483
6      462
7      143
4       33
8       16
3        6
Name: quality, dtype: int64
```

```
0      1
1079   1
1087   1
1086   1
1085   1
      ..
543    1
544    1
545    1
546    1
1597   1
Name: Id, Length: 1143, dtype: int64
```

- No junk present in our data

```
#data split
from sklearn.model_selection import train_test_split
np.random.seed(0)
df_train,df_test=train_test_split(data_wine,train_size=0.7,test_size=0
.3,random_state=100)
```

```
# Data Visualization
df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 800 entries, 97 to 792
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          800 non-null    float64
1   volatile acidity       800 non-null    float64
2   citric acid            800 non-null    float64
3   residual sugar         800 non-null    float64
```

```

4   chlorides          800 non-null    float64
5   free sulfur dioxide 800 non-null    float64
6   total sulfur dioxide 800 non-null    float64
7   density            800 non-null    float64
8   pH                 800 non-null    float64
9   sulphates          800 non-null    float64
10  alcohol            800 non-null    float64
11  quality            800 non-null    int64
12  Id                 800 non-null    int64

```

dtypes: float64(11), int64(2)

memory usage: 87.5 KB

df_train.shape

(800, 13)

df_test.info()

<class 'pandas.core.frame.DataFrame'>

Int64Index: 343 entries, 459 to 166

Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
0	fixed acidity	343 non-null	float64
1	volatile acidity	343 non-null	float64
2	citric acid	343 non-null	float64
3	residual sugar	343 non-null	float64
4	chlorides	343 non-null	float64
5	free sulfur dioxide	343 non-null	float64
6	total sulfur dioxide	343 non-null	float64
7	density	343 non-null	float64
8	pH	343 non-null	float64
9	sulphates	343 non-null	float64
10	alcohol	343 non-null	float64
11	quality	343 non-null	int64
12	Id	343 non-null	int64

dtypes: float64(11), int64(2)

memory usage: 37.5 KB

df_test.shape

(343, 13)

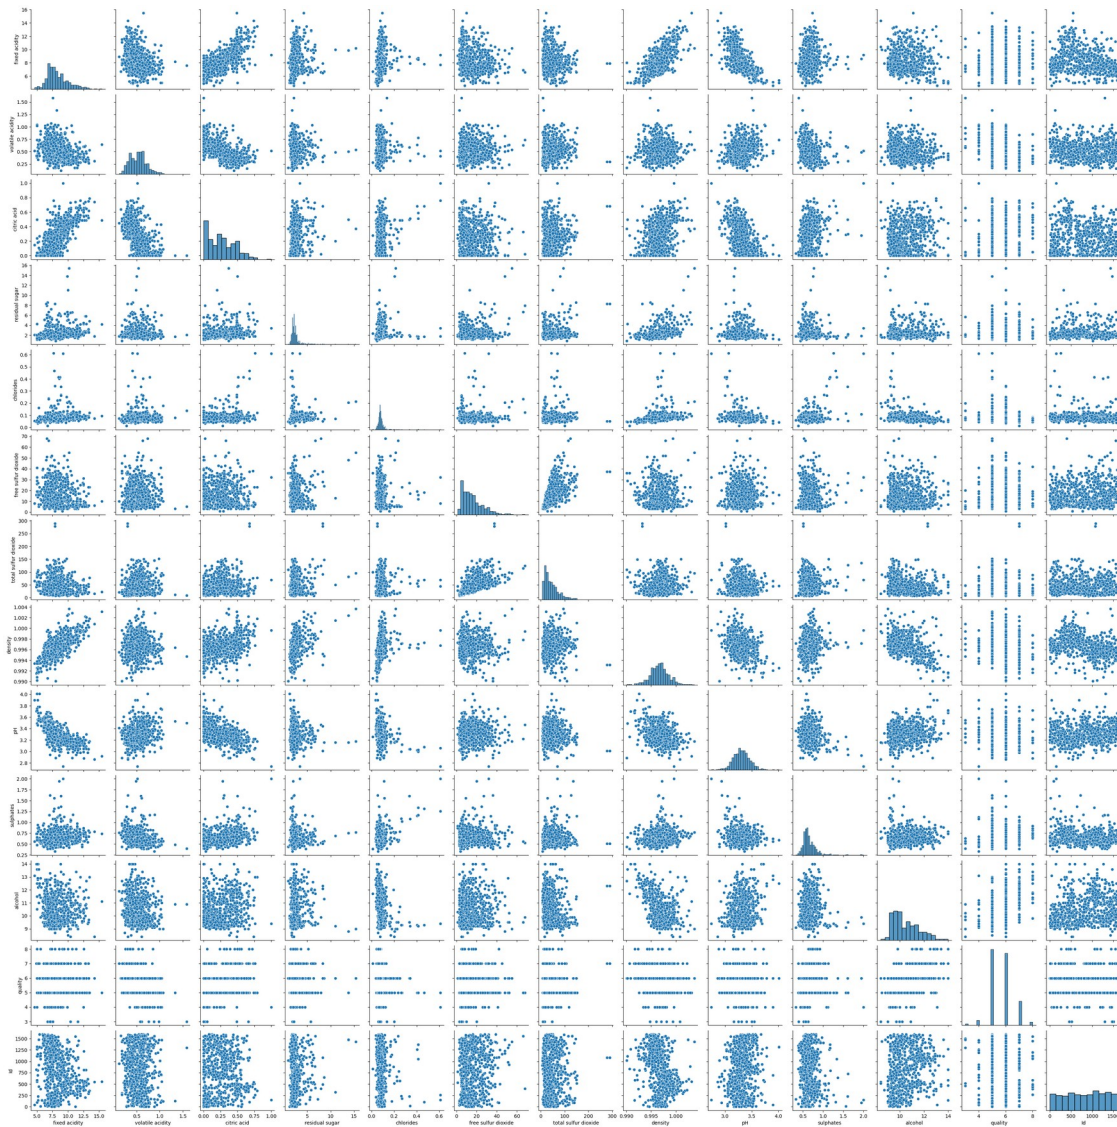
df_train.columns

```

Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual
sugar',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide',
'density',
      'pH', 'sulphates', 'alcohol', 'quality', 'Id'],
      dtype='object')

```


```
#using seabora lib
sns.pairplot(df_train)
plt.show()
```



Correlation plot

```
#using heatmap
```

```
plt.figure(figsize=(20,25))
sns.heatmap(data_wine.corr(), annot=True)
sns.light_palette("#a275ac", as_cmap=True)
plt.show()
```



Rescaling

```
from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler()
```

```
df_train.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar
chlorides \				
97	6.3	0.39	0.08	1.7
0.066				
353	8.7	0.69	0.31	3.0
0.086				
328	10.3	0.50	0.42	2.0
0.069				
191	6.9	0.54	0.04	3.0
0.077				
764	6.8	0.48	0.08	1.8
0.074				

	free sulfur dioxide	total sulfur dioxide	density	pH
sulphates \				
97	3.0	20.0	0.99540	3.34
0.58				
353	23.0	81.0	1.00020	3.48
0.74				
328	21.0	51.0	0.99820	3.16
0.72				
191	7.0	27.0	0.99870	3.69
0.91				
764	40.0	64.0	0.99529	3.12
0.49				

	alcohol	quality	Id
97	9.4	5	143
353	11.6	6	499
328	11.5	6	466
191	9.4	6	268
764	9.6	5	1085

```
df_train.columns
```

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual  
sugar',  
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide',  
      'density',  
      'pH', 'sulphates', 'alcohol', 'quality', 'Id'],  
      dtype='object')
```

```
df_train[:]=scaler.fit_transform(df_train[:])
```

```
df_train.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar
chlorides \				
97	0.155963	0.184932	0.08	0.055172

0.090150				
353	0.376147	0.390411	0.31	0.144828
0.123539				
328	0.522936	0.260274	0.42	0.075862
0.095159				
191	0.211009	0.287671	0.04	0.144828
0.108514				
764	0.201835	0.246575	0.08	0.062069
0.103506				

	free sulfur dioxide	total sulfur dioxide	density	pH
free sulfur dioxide				
97	0.029851	0.049470	0.391336	0.472441
0.149701				
353	0.328358	0.265018	0.743759	0.582677
0.245509				
328	0.298507	0.159011	0.596916	0.330709
0.233533				
191	0.089552	0.074205	0.633627	0.748031
0.347305				
764	0.582090	0.204947	0.383260	0.299213
0.095808				

	alcohol	quality	Id
97	0.178571	0.4	0.089543
353	0.571429	0.6	0.312461
328	0.553571	0.6	0.291797
191	0.178571	0.6	0.167815
764	0.214286	0.4	0.679399

```
df_train.columns
```

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide',
      'density',
      'pH', 'sulphates', 'alcohol', 'quality', 'Id'],
      dtype='object')
```

Model fitting

```
y_train=df_train.pop('quality')
```

```
X_train=df_train
```

```
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression
```

```
#rfe = recursive feature elimination
```

```
lm = LinearRegression()
lm.fit(X_train, y_train)
```

```
rfe = RFE(lm,n_features_to_select=9)
rfe = rfe.fit(X_train, y_train)

list(zip(X_train.columns,rfe.support_,rfe.ranking_))
```

```
[('fixed acidity', True, 1),
 ('volatile acidity', True, 1),
 ('citric acid', True, 1),
 ('residual sugar', True, 1),
 ('chlorides', True, 1),
 ('free sulfur dioxide', False, 3),
 ('total sulfur dioxide', False, 2),
 ('density', True, 1),
 ('pH', True, 1),
 ('sulphates', True, 1),
 ('alcohol', True, 1),
 ('Id', False, 4)]
```

Module Selection

```
col = X_train.columns[rfe.support_]
col
```

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual
sugar',
      'chlorides', 'density', 'pH', 'sulphates', 'alcohol'],
      dtype='object')
```

```
X_train.columns[~rfe.support_]
```

```
Index(['free sulfur dioxide', 'total sulfur dioxide', 'Id'],
      dtype='object')
```

```
X_train_rfe = X_train[col]
```

Module 1

#VIF check Variance Inflation Factor

```
from statsmodels.stats.outliers_influence import
variance_inflation_factor
vif = pd.DataFrame()
vif['Features'] = X_train_rfe.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe.values, i) for i
in range(X_train_rfe.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

	Features	VIF
5	density	1305.03
6	pH	979.43
8	alcohol	121.62
0	fixed acidity	71.94
7	sulphates	20.67


```

1 volatile acidity    16.42
2      citric acid    9.05
4      chlorides      6.58
3 residual sugar     5.04

import statsmodels.api as sm
X_train_lm1 = sm.add_constant(X_train_rfe)
lr1 = sm.OLS(y_train, X_train_lm1).fit()

```

#parameter

```

lr1.params

const                29.803076
fixed acidity         0.043829
volatile acidity     -1.204881
citric acid           -0.224131
residual sugar        0.023565
chlorides             -1.759261
density              -26.663378
pH                   -0.300273
sulphates             0.688950
alcohol              0.326336
dtype: float64

```

```
print(lr1.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          quality    R-squared:
0.395
Model:                  OLS        Adj. R-squared:
0.389
Method:                 Least Squares    F-statistic:
57.41
Date:                   Sat, 20 May 2023    Prob (F-statistic):
2.00e-80
Time:                   16:49:20    Log-Likelihood:
-785.83
No. Observations:      800    AIC:
1592.
Df Residuals:          790    BIC:
1638.
Df Model:              9

Covariance Type:       nonrobust

```

```

=====
=====

```

	coef	std err	t	P> t
--	------	---------	---	------

[0.025 0.975]

```
-----
-----
const          29.8031    29.870    0.998    0.319    -
28.830      88.437
fixed acidity   0.0438    0.037    1.183    0.237    -
0.029      0.117
volatile acidity -1.2049    0.166   -7.270    0.000    -
1.530     -0.880
citric acid    -0.2241    0.211   -1.061    0.289    -
0.639      0.190
residual sugar  0.0236    0.023    1.040    0.298    -
0.021      0.068
chlorides      -1.7593    0.597   -2.945    0.003    -
2.932     -0.587
density        -26.6634   30.485   -0.875    0.382    -
86.505     33.178
pH             -0.3003    0.262   -1.144    0.253    -
0.815      0.215
sulphates       0.6889    0.160    4.318    0.000
0.376      1.002
alcohol         0.3263    0.038    8.695    0.000
0.253      0.400
=====
```

```
=====
Omnibus:                7.025    Durbin-Watson:
1.999
Prob(Omnibus):          0.030    Jarque-Bera (JB):
9.664
Skew:                   -0.027    Prob(JB):
0.00797
Kurtosis:               3.536    Cond. No.
2.62e+04
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.62e+04. This might indicate that there are strong multicollinearity or other numerical problems.

module 2

```
X_train_new = X_train_rfe.drop(["residual sugar"], axis = 1)
from statsmodels.stats.outliers_influence import
variance_inflation_factor
vif = pd.DataFrame()
vif['Features'] = X_train_new.columns
vif['VIF'] = [variance_inflation_factor(X_train_new.values, i) for i
```

```

in range(X_train_new.shape[1]))
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif

```

	Features	VIF
4	density	1304.03
5	pH	978.89
7	alcohol	120.69
0	fixed acidity	71.88
6	sulphates	20.54
1	volatile acidity	16.23
2	citric acid	8.90
3	chlorides	6.55

```
X_train_lm2 = sm.add_constant(X_train_new)
```

```
lr2 = sm.OLS(y_train, X_train_lm2).fit()
```

```
lr2.params
```

const	11.550422
fixed acidity	0.026665
volatile acidity	-1.201916
citric acid	-0.200228
chlorides	-1.744789
density	-8.022114
pH	-0.394655
sulphates	0.651506
alcohol	0.345695
dtype: float64	

```
print(lr2.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          quality    R-squared:
0.395
Model:                  OLS        Adj. R-squared:
0.388
Method:                 Least Squares    F-statistic:
64.44
Date:                   Sat, 20 May 2023    Prob (F-statistic):
4.11e-81
Time:                   16:58:36    Log-Likelihood:
-786.37
No. Observations:      800    AIC:
1591.
Df Residuals:          791    BIC:
1633.

```

Df Model: 8

Covariance Type: nonrobust

```
=====
```

		coef	std err	t	P> t	
[0.025 0.975]						

const		11.5504	24.176	0.478	0.633	-
35.907	59.008					
fixed acidity		0.0267	0.033	0.804	0.422	-
0.038	0.092					
volatile acidity		-1.2019	0.166	-7.252	0.000	-
1.527	-0.877					
citric acid		-0.2002	0.210	-0.954	0.341	-
0.612	0.212					
chlorides		-1.7448	0.597	-2.922	0.004	-
2.917	-0.572					
density		-8.0221	24.666	-0.325	0.745	-
56.440	40.396					
pH		-0.3947	0.246	-1.602	0.109	-
0.878	0.089					
sulphates		0.6515	0.155	4.191	0.000	
0.346	0.957					
alcohol		0.3457	0.033	10.605	0.000	
0.282	0.410					

```
=====
```

Omnibus:	6.658	Durbin-Watson:
2.007		
Prob(Omnibus):	0.036	Jarque-Bera (JB):
9.003		
Skew:	-0.028	Prob(JB):
0.0111		
Kurtosis:	3.517	Cond. No.
2.09e+04		

```
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.09e+04. This might indicate that there are strong multicollinearity or other numerical problems.

module 3

X_train_new = X_train_rfe.drop(["density"], axis = 1)

```

from statsmodels.stats.outliers_influence import
variance_inflation_factor
vif = pd.DataFrame()
vif['Features'] = X_train_new.columns
vif['VIF'] = [variance_inflation_factor(X_train_new.values, i) for i
in range(X_train_new.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif

```

	Features	VIF
5	pH	161.56
7	alcohol	120.85
0	fixed acidity	41.09
6	sulphates	20.46
1	volatile acidity	16.35
2	citric acid	9.05
4	chlorides	6.19
3	residual sugar	5.03

```
X_train_lm3 = sm.add_constant(X_train_new)
```

```
lr3 = sm.OLS(y_train, X_train_lm).fit()
```

```
lr3.params
```

const	3.687835
fixed acidity	0.018087
volatile acidity	-1.222027
citric acid	-0.222621
residual sugar	0.011922
chlorides	-1.790639
pH	-0.433890
sulphates	0.649692
alcohol	0.351453

dtype: float64

```
print(lr3.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          quality    R-squared:
0.395
Model:                  OLS        Adj. R-squared:
0.389
Method:                 Least Squares    F-statistic:
64.51
Date:                   Sat, 20 May 2023    Prob (F-statistic):
3.51e-81
Time:                   20:55:48    Log-Likelihood:

```

-786.21
 No. Observations: 800 AIC:
 1590.
 Df Residuals: 791 BIC:
 1633.
 Df Model: 8

Covariance Type: nonrobust

[0.025 0.975]		coef	std err	t	P> t	

const		3.6878	0.816	4.517	0.000	
2.085	5.290					
fixed acidity		0.0181	0.023	0.803	0.422	-
0.026	0.062					
volatile acidity		-1.2220	0.165	-7.426	0.000	-
1.545	-0.899					
citric acid		-0.2226	0.211	-1.054	0.292	-
0.637	0.192					
residual sugar		0.0119	0.018	0.651	0.515	-
0.024	0.048					
chlorides		-1.7906	0.596	-3.003	0.003	-
2.961	-0.620					
pH		-0.4339	0.213	-2.033	0.042	-
0.853	-0.015					
sulphates		0.6497	0.153	4.244	0.000	
0.349	0.950					
alcohol		0.3515	0.024	14.545	0.000	
0.304	0.399					
=====						
=====						
Omnibus:		7.064		Durbin-Watson:		
2.002						
Prob(Omnibus):		0.029		Jarque-Bera (JB):		
9.732						
Skew:		-0.028		Prob(JB):		
0.00770						
Kurtosis:		3.537		Cond. No.		
529.						
=====						
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
# module 4
X_train_new = X_train_rfe.drop(["pH"], axis = 1)
from statsmodels.stats.outliers_influence import
variance_inflation_factor
vif = pd.DataFrame()
vif['Features'] = X_train_new.columns
vif['VIF'] = [variance_inflation_factor(X_train_new.values, i) for i
in range(X_train_new.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

	Features	VIF
5	density	215.27
7	alcohol	112.89
0	fixed acidity	53.44
6	sulphates	20.66
1	volatile acidity	16.33
2	citric acid	8.98
4	chlorides	6.37
3	residual sugar	5.03

```
X_train_lm4 = sm.add_constant(X_train_new)
```

```
lr4 = sm.OLS(y_train, X_train_lm4).fit()
```

```
lr4.params
```

const	48.994264
fixed acidity	0.073537
volatile acidity	-1.203757
citric acid	-0.206833
residual sugar	0.032520
chlorides	-1.637087
density	-46.963559
sulphates	0.720237
alcohol	0.301736
dtype:	float64

```
print(lr4.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable:          quality    R-squared:
0.394
Model:                  OLS        Adj. R-squared:
0.388
Method:                 Least Squares    F-statistic:
64.40
Date:                   Sat, 20 May 2023    Prob (F-statistic):
```

4.59e-81
Time: 20:58:12 Log-Likelihood:
-786.49
No. Observations: 800 AIC:
1591.
Df Residuals: 791 BIC:
1633.
Df Model: 8

Covariance Type: nonrobust

[0.025 0.975]		coef	std err	t	P> t	

const		48.9943	24.719	1.982	0.048	
0.471	97.518					
fixed acidity		0.0735	0.026	2.781	0.006	
0.022	0.125					
volatile acidity		-1.2038	0.166	-7.262	0.000	-
1.529	-0.878					
citric acid		-0.2068	0.211	-0.982	0.327	-
0.620	0.207					
residual sugar		0.0325	0.021	1.530	0.126	-
0.009	0.074					
chlorides		-1.6371	0.588	-2.785	0.005	-
2.791	-0.483					
density		-46.9636	24.794	-1.894	0.059	-
95.633	1.706					
sulphates		0.7202	0.157	4.581	0.000	
0.412	1.029					
alcohol		0.3017	0.031	9.806	0.000	
0.241	0.362					
=====						
=====						
Omnibus:		7.851	Durbin-Watson:			
2.001						
Prob(Omnibus):		0.020	Jarque-Bera (JB):			
11.188						
Skew:		-0.027	Prob(JB):			
0.00372						
Kurtosis:		3.577	Cond. No.			
2.09e+04						
=====						
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is

correctly specified.

[2] The condition number is large, 2.09e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Module 5

```
X_train_new = X_train_rfe.drop(["sulphates"], axis = 1)
from statsmodels.stats.outliers_influence import
variance_inflation_factor
vif = pd.DataFrame()
vif['Features'] = X_train_new.columns
vif['VIF'] = [variance_inflation_factor(X_train_new.values, i) for i
in range(X_train_new.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

	Features	VIF
5	density	1292.05
6	pH	979.09
7	alcohol	119.81
0	fixed acidity	71.93
1	volatile acidity	15.81
2	citric acid	9.02
4	chlorides	5.55
3	residual sugar	5.01

```
X_train_lm5 = sm.add_constant(X_train_new)
```

```
lr5 = sm.OLS(y_train, X_train_lm5).fit()
```

```
lr5.params
```

const	-6.149806
fixed acidity	0.007604
volatile acidity	-1.359458
citric acid	-0.166521
residual sugar	0.001509
chlorides	-0.819417
density	10.361338
pH	-0.494495
alcohol	0.373168
dtype:	float64

```
print(lr5.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable:          quality    R-squared:
0.381
```

Model:	OLS	Adj. R-squared:
0.375		
Method:	Least Squares	F-statistic:
60.90		
Date:	Sat, 20 May 2023	Prob (F-statistic):
2.18e-77		
Time:	21:00:45	Log-Likelihood:
-795.16		
No. Observations:	800	AIC:
1608.		
Df Residuals:	791	BIC:
1650.		
Df Model:	8	

Covariance Type: nonrobust

		coef	std err	t	P> t	
[0.025 0.975]						

const		-6.1498	29.004	-0.212	0.832	-
63.083	50.784					
fixed acidity		0.0076	0.036	0.208	0.835	-
0.064	0.079					
volatile acidity		-1.3595	0.164	-8.308	0.000	-
1.681	-1.038					
citric acid		-0.1665	0.213	-0.781	0.435	-
0.585	0.252					
residual sugar		0.0015	0.022	0.068	0.946	-
0.042	0.045					
chlorides		-0.8194	0.562	-1.457	0.146	-
1.923	0.285					
density		10.3613	29.579	0.350	0.726	-
47.701	68.424					
pH		-0.4945	0.261	-1.891	0.059	-
1.008	0.019					
alcohol		0.3732	0.036	10.272	0.000	
0.302	0.444					

Omnibus:	6.231	Durbin-Watson:
2.025		
Prob(Omnibus):	0.044	Jarque-Bera (JB):
8.337		
Skew:	-0.006	Prob(JB):
0.0155		
Kurtosis:	3.500	Cond. No.
2.52e+04		

```
=====
```

```
Notes:
```

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.52e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
# Module 6
```

```
X_train_new = X_train_rfe.drop(["chlorides"], axis = 1)
```

```
from statsmodels.stats.outliers_influence import
```

```
variance_inflation_factor
```

```
vif = pd.DataFrame()
```

```
vif['Features'] = X_train_new.columns
```

```
vif['VIF'] = [variance_inflation_factor(X_train_new.values, i) for i
```

```
in range(X_train_new.shape[1])]
```

```
vif['VIF'] = round(vif['VIF'], 2)
```

```
vif = vif.sort_values(by = "VIF", ascending = False)
```

```
vif
```

	Features	VIF
4	density	1227.36
5	pH	947.68
7	alcohol	114.61
0	fixed acidity	68.01
6	sulphates	17.43
1	volatile acidity	15.06
2	citric acid	8.36
3	residual sugar	5.01

```
X_train_lm6 = sm.add_constant(X_train_new)
```

```
lr6 = sm.OLS(y_train, X_train_lm6).fit()
```

```
lr6.params
```

const	34.502282
fixed acidity	0.064234
volatile acidity	-1.340852
citric acid	-0.395668
residual sugar	0.022011
density	-32.055519
pH	-0.162084
sulphates	0.517704
alcohol	0.338818

```
dtype: float64
```

```
print(lr6.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          quality    R-squared:
0.389
Model:                  OLS       Adj. R-squared:
0.383
Method:                 Least Squares    F-statistic:
62.89
Date:                   Sat, 20 May 2023    Prob (F-statistic):
1.72e-79
Time:                   21:02:48    Log-Likelihood:
-790.19
No. Observations:      800    AIC:
1598.
Df Residuals:          791    BIC:
1641.
Df Model:               8

Covariance Type:       nonrobust

```

```

=====
=====

```

		coef	std err	t	P> t	
[0.025	0.975]					

const		34.5023	29.971	1.151	0.250	-
24.331	93.335					
fixed acidity		0.0642	0.037	1.756	0.079	-
0.008	0.136					
volatile acidity		-1.3409	0.160	-8.383	0.000	-
1.655	-1.027					
citric acid		-0.3957	0.204	-1.940	0.053	-
0.796	0.005					
residual sugar		0.0220	0.023	0.967	0.334	-
0.023	0.067					
density		-32.0555	30.578	-1.048	0.295	-
92.078	27.967					
pH		-0.1621	0.259	-0.625	0.532	-
0.671	0.347					
sulphates		0.5177	0.149	3.467	0.001	
0.225	0.811					
alcohol		0.3388	0.037	9.041	0.000	
0.265	0.412					

```

=====
=====
Omnibus:               8.039    Durbin-Watson:
2.021

```

Prob(Omnibus):	0.018	Jarque-Bera (JB):
11.633		
Skew:	-0.005	Prob(JB):
0.00298		
Kurtosis:	3.591	Cond. No.
2.62e+04		

=====

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.62e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Suggestions are always welcome.

Thank You !